

AI: definitions

Strong AI: general problem solver, can solve many problems, very dynamic Weak AI: solves **fewer** problems (usually 1, but not always)

1. percepts: input, e.g. what is seen
2. sensors: receive/parse input
3. functions: process input, produce execution steps
4. actuators: performs actions
5. actions: output env --percepts--> sensors > functions > actuators --actions--> env (agents = sensors + functions + actuators) Rational agent: Given [percept sequence, prior knowledge, set of actions, performance measure], optimises performance measure

Agent function: P, percept sequence --agent function f--> action A

Environment properties:

1. fully observable(can see entire board) vs partially observable(can only see part of board)
2. deterministic(same actions same outcome) vs stochastic(gamba)
3. episodic(next action will not affect future choices, **can plan**) vs sequential(next action e.g. pacman move character, will affect future decisions)
 1. If we can plan, e.g. produce sequence of actions (then execute sequentially e.g. sudoku), considered episodic.
4. Discrete(fixed set of states possible) vs continuous (infinite number of states possible)
5. Single agent (play with urself) vs Multi agent(against another player/AI)
6. Known (know mechanics, performance metrics etc) vs Unknown (black box)
7. Static vs Dynamic(Environment changes while agent is deciding on action)

Types of agents:

1. Reflex agent: **directly maps** percepts to actions (e.g. if at 0,0 go left; if at 0,1 go up; ...)
2. Model based reflex action: generalizes percepts into models (e.g. groups similar situations into same case, produce action)
3. Goal based agent: given state and actions, definition of goal, tries to satisfy goal
4. Utility based agent: like goal, but **optimizes** utility, e.g. chess, takes best move (more complex than goal, need to understand utility not just goal)

Search space definition

1. State representation si: data describing instance i of the environment, e.g. current position, walls etc.
2. Goal test: $\text{isGoal}(s_i) \Rightarrow \text{boolean}$
3. Actions: $\text{actions}(s_i) \Rightarrow \text{Action}[]$, possible actions at state i
4. Action costs: $\text{cost}(s_i, a_j, s_k) \Rightarrow \text{number}$, cost of taking an action a_j to transition s_i to s_k , generally ≥ 0 .
5. Transition model: $T(s_i, a_j) \Rightarrow s_k$, result of executing a_j at s_i . summarized as a graph problem, can find traversal start state to end state

General search algorithm

Uninformed search:

```
const frontier = state.edges
while (frontier) {
  const current = frontier.pop()
  if(isGoal(current.state)) return current.getpath()
  for a in actions(current.state) frontier.push(Node(T(state.current, a)))
}
```

- Node = state + parent + actions + cost + depth. e.g. current path A -> B -> C, Node(C, parent = [A, B], action = 'moveup', ...).
- uninformed search = no domain knowledge beyond search problem formulation
- typical uninformed search algo differ in frontier data struct: BFS = queue, Uniform Cost search = prio queue, Depth first = stack etc.
- Correctness: algorithm is **Complete** if solution will be found when one exists, raise error if no solution found. **Optimal** = find solution with lowest cost
- by default **late goal test** (test goal on neighbours before adding to frontier) except for BFS Breadth First Search (default **early goal test**):

1. Queue frontier
2. $O(b^d)$ time/space complexity: b = branching factor, how many branch; d = depth of tree
3. "**Complete**" (on condition: finite branching factor && (finite search space OR has a solution)), Suboptimal (no

concern about cost)

4. Can be improved using Early goal test (test before push) instead of Late goal test (original) --> saves b^k+1 where k is the depth of the goal

Uniform Cost Search (Dijkstra)

1. Queue = prio queue w/ cost prio
2. Time/Space complexity $O(b^e)$ where $e=1 + \text{floor}(\text{Optimal path cost } C / k)$, k is a small positive constant
3. "**Complete**"(on same BFS condition), optimal (late goal) --> early goal may not be optimal w/ graph search

Depth First Search:

1. Time complexity $O(b^m)$, space = $O(bm)$, m == max depth --> bm because at max space req = max depth b^m no. of nodes in stack
2. Incomplete under BFS condition: if infinite search space, aka depth, dfs continue forever
3. suboptimal
4. Can be improved by backtracking instead of each node storing entire path back to node. $O(m)$ space
5. **NOTE:** pop in reverse order due to stack: pop A push B push C, pop C, pop B.

Depth limited/Iterative deepening

1. Depth limited = DFS with depth limit l, only search up to l depth, no actions at l depth
2. Depth limited same guarantees as DFS
3. Iterative deepening: increase depth limit by 1 each time until solution found
4. Iterative deepening guarantees completeness under BFS conditions, but with space complexity of dfs $O(bm)$. repeated calculation of smaller depths overhead ~11%, insignificant

Trees vs graphs

1. v1: (before pushing to frontier) if child not in reached: frontier.push(child) && reached[child.state] = child;
 2. v2: (before pushing to frontier) if child not in reached or child.cost < reached[child.state].cost: frontier.push(child) && reached[child.state] = child;
 3. v3: when exploring node: reached[node.state] = node; before push: if child not in reached: frontier.push(child); (like "late reaching")
- Graphs can have cycles; to prevent revisits, can add visited hashmap (Graph search v1), check map before pushing adjacent nodes; Graph search v2: if revisit has lower cost update cost and add to frontier (revisit it)
 - Graph search **default graph search v1**, all BFS UCS DFS DLS IDS space time $O(|V| + |E|)$, no check cheaper path

Informed search:

General Best First Algo:

```
while len(frontier) > 0:
  node = frontier.pop()
  if is_goal(node.state): # Late goal test
    return node
# graph search v2
for child in [Node(state=apply(node.state, action), parent=node, cost=node.cost + action.cost)
  if child.state in reached or child.cost < reached[child.state].cost:
    reached[child.state] = child
  frontier.append(child)
```

Greedy Best First Search:

1. UCS except prio determined by $h(n)$ heuristic: distance from g.
2. sub optimal: when adding child nodes to frontier, does not consider distance from parent to child, only sorted by distance child to goal
3. (tree search) incomplete: if A --> B and B --> A, both A/B same h, can end up going back and forth b/w A and B
4. (graph search) complete if finite state space: will visit entire space

A* search:

1. Greedy Best First + account for alr incurred costs
2. $f(n) = g(n) + h(n) \rightarrow g(n)$ = actual path cost, $h(n)$ = estimated cheapest path to G; $h(n)$ gets more accurate the closer it is to G
3. completeness, same criteria as ucs
4. optimality: optimal if $h(n)$ is admissible (tree, graph v2 late goal); early goal/ graph v1 may skip optimal
5. if $h(n)$ consistent, graph search optimal by v3 (v3 fixes v1 --> v1 cannot add more than 1 path to G since G marked as reached too early)
6. note: if $h_1(n) \geq h_2(n)$, then h_1 *dominates* h_2 , aka h_1 more efficient, higher accuracy, need to try fewer paths

(assuming both admissible)

admissible heuristics:

- $h(n)$ considered admissible if it never overestimates cost (to goal): $h(n) \leq h^*(n)$, paths ending at goal exact, paths not ending at goal overestimated
- e.g. euclidean $\sqrt{a^2 + b^2}$ always underestimates

consistent heuristics:

- requirement: $h(\text{parent}) \leq h(\text{child}) + \text{cost}(\text{parent}, \text{child})$, so f costs can be monotonically increasing
- if consistent, then also admissible

Effective branching factor

- More efficient heuristics, less number of paths with cost \leq optimal path, less nodes explored
- can be evaluated with estimated branching factor: N nodes explored, solution at d depth, find b s.t. $N + 1 = (b^{d+1} - 1) / (b - 1)$, compare b
- derived from $N + 1 = \sum(b^i)[1 \rightarrow d] \rightarrow Gp$ formula $(a^n - 1)/(n - 1)$, $a = b$, $n = d$.

Relaxing problems: to generate a heuristic for a problem, sometimes we can simplify the problem, e.g. path finding, every cost is 1

Gradient descent

- local search: greedy approach, maintain best successor. good for large/infinite search spaces, but might not guarantee goal
- hill climb algorithm:

```
while True: # Like f(n) = -h(n)
    neighbour = max(current.successors) # steepest hill/ greedy
    if value(neighbour) <= value(current):
        return current
    current = neighbour
    if is_goal(current):
        return current
```

- complete state formulation: **each state is a potential "solution"**. either guess and check neighbours or move to states with higher value $f(n) = -h(n)$.
 - good for problems where "path" to goal not important
 - $-h(n)$ simply to satisfy hill climb name, otherwise $h(n)$ descending, lower better
- sideways move: instead of $\text{neighbour} < \text{current}$ use $\text{neighbour} \leq \text{current}$, so can traverse plateaus
- stochastic hill climb: choose randomly among states better than current (not just best). takes longer, but may find better solution (Stochastic gradient descent = pick random sets of training data for each iteration of learning in ML)
- First choice hill climb: when too many successors, randomly generate successor until 1 better value (instead of generate all)
- random restart: add outer loop to pick a new random starting state, reset until found
 - expected number of steps: $x + (1-p)/p * y$, $x = E(\text{steps to find actual goal})$, $y = E(\text{steps to find local maxima})$, $p = \text{probability to find global maxima}$, $(1-p)/p = E(\text{no. failures})$
- local beam search: store k states instead of current state, each iteration generate successors for all k , repeat for best k (or random out of better k a la stochastic hill climb). goal test for each k after selecting k

Constraint satisfaction Problems

- systematic approach, do not search states where constraint not satisfied. prune invalid subtrees
- Only consider unary/binary constraints, ignore global constraints unless its CSP formulation question
- constraints can be dynamic, e.g. simultaneous equations, each eqn = constraint
 - notation: $\langle (\text{scope}), \text{rel} \rangle$: e.g. $\langle (x_1, x_2), x_1 > x_2 \rangle$. |scope|, single = unary, double = binary, >2 global
 - hypergraphs: link variables (circles/vertexes) via relationships (edges). linking vertex (box) can also be used to join global constraints

```
# CSP (DFS) general algorithm
while is_incomplete(assignments):
    try: # if nothing to assign return failure
        current = assign_to_non_assigned(current)
        if current.consistent():
            assignments.push(current)
    except:
        return failure
return assignments
```

- use DFS traverse. keep assigning until fail
- can backtrack if fail (like DFS) (or look ahead check if fail before assigning)
- state representation: initial state all variable unassigned. domain = possible values, variables = stuff that need to be assigned a value
- action: assign values to variables. no costs/evaluation function used
- goal test: all constraints satisfied
- CSP search tree leaves = $n!m^n \rightarrow$ first level nm leaf states, subtree of each = $(n-1)^m, (n-2)^m \dots$

```
def backtrack(csp, assignment):
    if assignment.complete():
        return assignment
    var = select_unassigned_var(csp, assignment)
    for value in domain.sort(get_value_sorter(csp, assignment, var)):
        if consistent(value, var, assignment):
            assign(value, var, assignment)
            if infer(assignment) != failure:
                csp.append(infer(assignment))
                result = backtrack(csp, assignment)
                if result != failure:
                    return result
            unassign(value, var, assignment) # backtrack
    return failure
```

- constraint graph: simple vertex, circle: variable, linking vertex, square: global constraints, edge: links variables in scope of constraint
 - unary constraints = node itself, binary constraint = edge
- variable order heuristics: (select_unassigned_var)
 1. Minimum Remaining values (MRV): Choose variable with fewest consistent values (e.g. map colouring, area on map with fewest possible colours) \rightarrow finds inconsistent subtree fast, prunes them
 2. Degree heuristic: used for tie breaking stuff like MRV. Select variable w/ most constraints among unassigned variables (e.g. map colouring, area adjacent to most other areas)
- value order heuristics: (get_value_sorter)
 1. Least Constraining value (LCV): after choosing a variable, choose value that rules out the fewest subsequent choices (maximize subsequent choices, prioritize flexibility of subsequent assignments)
- inferences infer(assignment):
 - forward checking: check remaining legal values for unassigned variables, terminate if any has no legal values (does not provide early detection for all failures, only 1 step ahead)
 - constraint propagation: solves forward checking flaws, traverses entire constraint graph to make sure still consistent (note global constraints can be converted to binary from invisible node to all vertexes)
 - node consistent: unary constraints satisfied (vertexes) **unary constraint checked as preprocessing before backtracking algo**
 - arc consistent: binary constraints satisfied (edges). variable X_i is arc consistent wrt X_j iff for all values D_i there exists some value D_j that satisfies binary constraint. Note arc == directed, binary constraint = 2 arcs

```
def arc_consistent(csp): #aka AC-3 Algo O(n^2 d^3), d = domain size
    arcs = csp.get_arcs()
    while not arcs.empty():
        arc = arcs.pop()

    # revise domains
    revised = False
    for x in arc.x_i.domain:
        # no value in Xj domain can satisfy constraint with x across arc
        if not can_satisfy_constraint(csp, x, arc.x_j.domain):
            # x cannot be used in arc, does not satisfy constraint
            delete(x)
            revised = True

    if revised:
        if len(arc.x_i.domain) < 1:
            return False
        for neighbour in arc.x_i.neighbours:
            arcs.append((x_i, neighbour))
    return True
```